

ML-HW3

Xiaoman Xu

2023-02-08

```
library('splines')      ## for 'bs'
library('dplyr')        ## for 'select', 'filter', and others
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library('magrittr')     ## for '%<>%' operator
#install.packages('glmnet')
library('glmnet')        ## for 'glmnet'
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-6
```

1. Use the prostate cancer data.

```
## load prostate data
prostate <-
  read.table(url(
    'https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data'))
```

2. Use the cor function to reproduce the correlations listed in HTF Table 3.1, page 50.

```
cor(prostate)
```

```
##          lcavol      lweight      age      lbph      svi
## lcavol    1.00000000  0.280521380  0.2249999  0.027349703  0.53884500
## lweight   0.28052138  1.000000000  0.3479691  0.442264399  0.15538490
## age       0.22499988  0.347969112  1.0000000  0.350185896  0.11765804
## lbph      0.02734970  0.442264399  0.3501859  1.000000000 -0.08584324
## svi       0.53884500  0.155384903  0.1176580 -0.085843238  1.00000000
## lcp       0.67531048  0.164537142  0.1276678 -0.006999431  0.67311118
## gleason   0.43241706  0.056882093  0.2688916  0.077820447  0.32041222
## pgg45     0.43365225  0.107353785  0.2761124  0.078460018  0.45764762
## lpsa      0.73446033  0.433319382  0.1695928  0.179809404  0.56621822
## train     -0.04654347 -0.009940658  0.1776155 -0.029939957  0.02679950
##          lcp      gleason      pgg45      lpsa      train
## lcavol    0.675310484  0.43241706  0.43365225  0.73446033 -0.046543468
## lweight   0.164537142  0.05688209  0.10735379  0.43331938 -0.009940658
## age       0.127667752  0.26889160  0.27611245  0.16959284  0.177615517
## lbph      -0.006999431  0.07782045  0.07846002  0.17980940 -0.029939957
## svi       0.673111185  0.32041222  0.45764762  0.56621822  0.026799505
## lcp       1.000000000  0.51483006  0.63152825  0.54881317 -0.037427296
## gleason   0.514830063  1.00000000  0.75190451  0.36898681 -0.044171456
## pgg45     0.631528246  0.75190451  1.00000000  0.42231586  0.100516371
## lpsa      0.548813175  0.36898681  0.42231586  1.00000000 -0.033889743
## train     -0.037427296 -0.04417146  0.10051637 -0.03388974  1.000000000
```

```
# heatmap(cor(prostate), Rowv=NA, Colv=NA, symm=T,
#          main='Correlation among inputs' )
```

3. Treat lcavol as the outcome, and use all other variables in the data set as predictors. & 4. With the training subset of the prostate data, train a least-squares regression model with all predictors using the lm function.

```
## split prostate into testing and training subsets
prostate_train <- prostate %>%
  filter(train == TRUE) %>%
  select(-train)

summary(prostate_train)
```

```
##          lcavol      lweight      age      lbph
## Min.      :-1.3471  Min.      :2.375  Min.      :41.00  Min.      :-1.38629
## 1st Qu.:  0.4883  1st Qu.: 3.330  1st Qu.: 61.00  1st Qu.: -1.38629
## Median :  1.4679  Median : 3.599  Median : 65.00  Median : -0.05129
## Mean      : 1.3135  Mean      : 3.626  Mean      : 64.75  Mean      : 0.07144
## 3rd Qu.:  2.3491  3rd Qu.: 3.884  3rd Qu.: 69.00  3rd Qu.:  1.54751
## Max.      :  3.8210  Max.      : 4.780  Max.      : 79.00  Max.      :  2.32630
```

```
##          svi          lcp          gleason          pgg45
## Min.    :0.0000   Min.    :-1.3863   Min.    :6.000   Min.    : 0.00
## 1st Qu.:0.0000   1st Qu.: -1.3863   1st Qu.:6.000   1st Qu.: 0.00
## Median :0.0000   Median :-0.7985   Median :7.000   Median : 15.00
## Mean    :0.2239   Mean    :-0.2142   Mean    :6.731   Mean    : 26.27
## 3rd Qu.:0.0000   3rd Qu.: 0.9948   3rd Qu.:7.000   3rd Qu.: 50.00
## Max.    :1.0000   Max.    : 2.6568   Max.    :9.000   Max.    :100.00
##          lpsa
## Min.    :-0.4308
## 1st Qu.: 1.6673
## Median : 2.5688
## Mean    : 2.4523
## 3rd Qu.: 3.3652
## Max.    : 5.4775
```

```
prostate_test <- prostate %>%
  filter(train == FALSE) %>%
  select(-train)
```

```
## predict lcavol consider all other predictors
## lm fits using L2 loss
fit <- lm(lcavol ~ ., data=prostate_train)
summary(fit)
```

```
##
## Call:
## lm(formula = lcavol ~ ., data = prostate_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.71027 -0.50138  0.03103  0.51352  1.35376
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.173357   1.526699  -1.424   0.160
## lweight      -0.113370   0.236639  -0.479   0.634
## age           0.020102   0.013571   1.481   0.144
## lbph         -0.056981   0.072525  -0.786   0.435
## svi           0.035116   0.313526   0.112   0.911
## lcp           0.418455   0.099521   4.205 9.16e-05 ***
## gleason      0.224387   0.198812   1.129   0.264
## pgg45        -0.009113   0.005451  -1.672   0.100 .
## lpsa          0.575455   0.107235   5.366 1.47e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7116 on 58 degrees of freedom
## Multiple R-squared:  0.7118, Adjusted R-squared:  0.672
## F-statistic: 17.9 on 8 and 58 DF, p-value: 3.999e-13
```

```
coef(fit)
```

```
## (Intercept)      lweight          age          lbph          svi          lcp
```

```
## -2.17335699 -0.11336968 0.02010188 -0.05698125 0.03511645 0.41845469
##      gleason      pgg45      lpsa
## 0.22438690 -0.00911273 0.57545508
```

```
residuals(fit)
```

```
##      1      2      3      4      5      6
## 0.30510051 -0.36214995 -0.31372527 -0.57604668 1.00855250 -0.80121337
##      8     11     12     13     14     16
## 0.93154303 -0.04430734 -1.45525306 1.00784627 0.84369636 1.05697269
##     17     18     19     20     21     23
## -1.07296621 1.10894970 -0.58433985 -0.20987943 0.73401898 -0.96212934
##     24     27     29     30     31     33
## 0.60155186 0.08773062 0.30844179 0.46587959 -0.39203706 0.51616536
##     35     37     38     39     40     41
## -1.00794601 -1.08327905 -0.58060313 0.68114682 -0.52751225 -0.16747733
##     43     45     46     47     51     52
## -0.46010692 0.15965010 0.70751569 0.13993772 0.20829842 0.83814762
##     56     58     59     60     61     63
## 0.30174358 -0.21418255 -0.61609540 0.27427103 -0.59282089 0.37445320
##     67     68     69     70     71     72
## 0.20410928 0.51088167 -1.71027355 -0.30736902 -0.18635786 -0.24666610
##     75     76     77     78     79     81
## 0.03553988 0.23104706 0.17458591 0.89281112 -0.27734203 -0.62839982
##     82     83     85     86     87     88
## -0.06963027 0.03103464 -0.65441115 0.38724844 0.66414753 -0.63266711
##     89     90     91     92     93     94
## -0.25266290 -0.87963313 1.35375660 0.70838106 -0.13844098 0.62802337
##     96
## -0.47525498
```

5. Use the testing subset to compute the test error (average squared-error loss) using the fitted least-squares regression model.

```
## functions to compute testing/training error w/lm
L2_loss <- function(y, yhat)
  (y-yhat)^2
error <- function(dat, fit, loss=L2_loss)
  mean(loss(dat$lcavol, predict(fit, newdata=dat)))

## testing error
error(prostate_test, fit)
```

```
## [1] 0.5084068
```

6. Train a ridge regression model using the `glmnet` function, and tune the value of `lambda` (i.e., use guess and check to find the value of `lambda` that approximately minimizes the test error).

```
## use glmnet to fit lasso
## glmnet fits using penalized L2 loss
## first create an input matrix and output vector
form <- lcavol ~ lweight + age + lbph + lcp + pgg45 + lpsa + svi + gleason
x_inp <- model.matrix(form, data=prostate_train)
y_out <- prostate_train$lcavol
fit <- glmnet(x=x_inp, y=y_out, lambda=seq(0.5, 0, -0.05))
print(fit$beta)
```

```
## 9 x 11 sparse Matrix of class "dgCMatrix"
```

```
##      [[ suppressing 11 column names 's0', 's1', 's2' ... ]]
```

```
##
## (Intercept) .          .          .          .          .          .
## lweight      .          .          .          .          .          .
## age          .          .          .          .          .          .
## lbph         .          .          .          .          .          .
## lcp          0.1473018 0.1714414 0.1955919 0.2197423 0.2438928 0.2680433
## pgg45        .          .          .          .          .          .
## lpsa         0.2535992 0.2816134 0.3096214 0.3376294 0.3656374 0.3936454
## svi          .          .          .          .          .          .
## gleason      .          .          .          .          .          .
##
## (Intercept) .          .          .          .          .
## lweight      .          .          .          .          -0.113959029
## age          .          0.0006179005 0.005823836 0.010987417 0.020114429
## lbph         .          .          .          -0.004060083 -0.056959719
## lcp          0.2921937 0.3160341041 0.337951116 0.352226351 0.418442122
## pgg45        .          .          .          .          -0.009110807
## lpsa         0.4216534 0.4489636282 0.470877681 0.491246304 0.575583816
## svi          .          .          .          0.026126960 0.035184640
## gleason      .          .          .          0.011159844 0.224210312
```

```
## functions to compute testing/training error with glmnet
error <- function(dat, fit, lam, form, loss=L2_loss) {
  x_inp <- model.matrix(form, data=dat)
  y_out <- dat$lcavol
  y_hat <- predict(fit, newx=x_inp, s=lam) ## see predict.elnet
  mean(loss(y_out, y_hat))
}
```

```
## train_error at lambda=0
error(prostate_train, fit, lam=0, form=form)
```

```
## [1] 0.4383709
```

```
## testing error at lambda=0
error(prostate_test, fit, lam=0, form=form)
```

```
## [1] 0.5084581
```

```
## train_error at lambda=0.03
error(prostate_train, fit, lam=0.05, form=form)
```

```
## [1] 0.4708637
```


```
## testing error at lambda=0.03
error(prostate_test, fit, lam=0.05, form=form)
```

```
## [1] 0.4945864
```

7. Create a figure that shows the training and test error associated with ridge regression as a function of lambda

```
## compute training and testing errors as function of lambda
err_train_1 <- sapply(fit$lambda, function(lam)
  error(prostate_train, fit, lam, form))
err_test_1 <- sapply(fit$lambda, function(lam)
  error(prostate_test, fit, lam, form))

## plot test/train error
plot(x=range(fit$lambda),
     y=range(c(err_train_1, err_test_1)),
     xlim=rev(range(fit$lambda)),
     type='n',
     xlab=expression(lambda),
     ylab='train/test error')
points(fit$lambda, err_train_1, pch=19, type='b', col='darkblue')
points(fit$lambda, err_test_1, pch=19, type='b', col='darkred')
legend('topright', c('train', 'test'), lty=1, pch=19,
      col=c('darkblue', 'darkred'), bty='n')
```



```
colnames(fit$beta) <- paste('lam =', fit$lambda)
print(fit$beta %>% as.matrix)
```

```
##          lam = 0.5 lam = 0.45 lam = 0.4 lam = 0.35 lam = 0.3 lam = 0.25
## (Intercept) 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## lweight     0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## age         0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## lbph        0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## lcp         0.1473018 0.1714414 0.1955919 0.2197423 0.2438928 0.2680433
## pgg45       0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## lpsa        0.2535992 0.2816134 0.3096214 0.3376294 0.3656374 0.3936454
## svi         0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## gleason     0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
##          lam = 0.2   lam = 0.15   lam = 0.1   lam = 0.05   lam = 0
## (Intercept) 0.0000000 0.0000000000 0.000000000 0.000000000 0.000000000
## lweight     0.0000000 0.0000000000 0.000000000 0.000000000 -0.113959029
## age         0.0000000 0.0006179005 0.005823836 0.010987417 0.020114429
## lbph        0.0000000 0.0000000000 0.000000000 -0.004060083 -0.056959719
## lcp         0.2921937 0.3160341041 0.337951116 0.352226351 0.418442122
## pgg45       0.0000000 0.0000000000 0.000000000 0.000000000 -0.009110807
## lpsa        0.4216534 0.4489636282 0.470877681 0.491246304 0.575583816
## svi         0.0000000 0.0000000000 0.000000000 0.026126960 0.035184640
## gleason     0.0000000 0.0000000000 0.000000000 0.011159844 0.224210312
```

8. Create a path diagram of the ridge regression analysis, similar to HTF Figure 3.8

```
## plot path diagram
plot(x=range(fit$lambda),
     y=range(as.matrix(fit$beta)),
     type='n',
     xlab=expression(lambda),
     ylab='Coefficients')
for(i in 1:nrow(fit$beta)) {
  points(x=fit$lambda, y=fit$beta[i,], pch=19, col='blue')
  lines(x=fit$lambda, y=fit$beta[i,], col='blue')
}
text(x=0, y=fit$beta[,ncol(fit$beta)],
     labels=rownames(fit$beta),
     xpd=NA, pos=4, srt=45)
abline(h=0, lty=3, lwd=2)
```

